

NoSQL vs SQL Database Comparison Report

Author: Benjamin Wuermli

Date: June 14, 2025

Course: Datenbankdesign und Implementierung (ddi)

Summary

This report compares MongoDB (NoSQL) and PostgreSQL (SQL) for an e-commerce product review system. With 10,000 products and 0 to 55 reviews per product.

Performance patterns, query complexities and flexibility show each database's strengths.

Key Finding: MongoDB excels at pre-aggregated data and simple lookups, while PostgreSQL is better for indexed filtering complex queries and situations where structured data is crucial.

Problem Description

Use Case: E-commerce product catalog with user reviews

Data Volume: - 10,000 products across 5 categories (Electronics, Books, Clothing, Home, Sports) - 283,887 user reviews (average 28 reviews per product) - Relationship: One-to-many (Product → Reviews)

Why NoSQL Makes Sense: - Products have varying attributes by category - Reviews naturally nest within products - Read-heavy workload (browsing > purchasing) - Flexible schema for different product types

Furthermore, MongoDB's schema-less nature should provide significant advantages in adapting to evolving data requirements. Unlike relational databases, where adding a new attribute requires altering the table schema, MongoDB allows you to simply add the new field to the relevant documents without affecting existing data. This flexibility is beneficial in this environment where product catalogs are constantly changing and new attributes are frequently introduced.

For example, if we wanted to add a "color" attribute to only the "Clothing" category, we could do so without modifying the structure of documents in other categories.

Conceptual Model

Relationship: One-to-Many (1:N)

- One PRODUCT can have many REVIEWS

- One REVIEW belongs to exactly one PRODUCT
- Foreign Key: REVIEW.product_id → PRODUCT.id

Further Rules: - Products must have unique IDs

- Reviews must reference existing products
- Ratings constrained to 1-5 scale
- Review must have dates

the shema

Database Implementations

PostgreSQL (SQL) Design

```
CREATE TABLE products (
    id SERIAL PRIMARY KEY,
    name VARCHAR(255) NOT NULL,
    category VARCHAR(100) NOT NULL,
    price DECIMAL(10,2) NOT NULL,
    description TEXT
);

CREATE TABLE reviews (
    id SERIAL PRIMARY KEY,
    product_id INTEGER REFERENCES products(id),
    user_name VARCHAR(100) NOT NULL,
    rating INTEGER CHECK (rating >= 1 AND rating <= 5),
    review_text TEXT,
    date TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

CREATE INDEX idx_reviews_product_id ON reviews(product_id);
CREATE INDEX idx_products_category ON products(category);
```

- **Normalization:** 3NF (separate tables)
- **Foreign Key:** reviews.product_id → products.id
- **Indexes:** On product_id, category

MongoDB (NoSQL) Design

```
{
  "_id": 1,
  "name": "Product Name",
  "category": "Electronics",
  "price": 299.99,
  "description": "Product description",
  "reviews": [
    {
      "id": 1,
      "user_name": "john_doe",
      "rating": 5,

```

```

        "review_text": "Great product!",
        "date": "2024-06-01T10:00:00Z"
    }
],
"avg_rating": 4.2,
"review_count": 15
}

```

- **Structure:** Denormalized documents
- **Embedded:** Reviews within products
- **Pre-calculated:** Average ratings and counts

Performance Results

the comparisons

Query Performance Comparison

Operation	MongoDB (ms)	PostgreSQL (ms)	Winner	Speedup
Get Average Ratings	0.0169	0.0806	MongoDB	4.77x
Get Product with Reviews	0.0006	0.0005	PostgreSQL	1.20x
Get Products by Category	0.0615	0.0022	PostgreSQL	27.95x
Get Products by Price Range	0.0081	0.0409	MongoDB	5.05x
Get Products with High Ratings	0.0071	0.0739	MongoDB	10.41x
Get Products with Keyword Reviews	0.3024	0.0432	PostgreSQL	7.00x

Query Complexity Analysis

MongoDB Advantages (Complexity = 1)

```
// Get product with all reviews - Single query
db.products.find({"_id": 1})
```

```
// Get average ratings
db.products.find({}, {"avg_rating": 1, "review_count": 1})
```

```
// Get products with keyword reviews
db.products.find({"reviews.review_text": {"$regex": keyword, "$options": "i"}}
```

PostgreSQL Complexity (Complexity = 2-3)

```
-- Get product with reviews - JOIN required
SELECT p.*, r.user_name, r.rating, r.review_text
FROM products p
LEFT JOIN reviews r ON p.id = r.product_id
WHERE p.id = 1;
```

```
-- Get average ratings - Complex aggregation
SELECT p.id, p.name, AVG(r.rating), COUNT(r.id)
```

```
FROM products p
LEFT JOIN reviews r ON p.id = r.product_id
GROUP BY p.id, p.name;
```

```
-- Get products with keyword reviews
SELECT p.*
FROM products p
WHERE EXISTS (
    SELECT 1
    FROM reviews r
    WHERE r.product_id = p.id
    AND r.review_text LIKE %s
-- with parameter "%" + keyword + "%"
```

Schema Flexibility

Adding New Product Attributes

PostgreSQL Process: 1. ALTER TABLE products ADD COLUMN brand VARCHAR(100); 2. ALTER TABLE products ADD COLUMN weight DECIMAL(10,2); 3. Update application code (Models) 4. Potentially many empty fields in the database

MongoDB Process: 1. Insert document with new fields - Done!

Winner: MongoDB - No schema migrations needed

Measurement Criteria Summary

1. Query Performance

- **Mixed Results:** Each database excels in different scenarios
- **MongoDB:** Better for embedded data access and pre-aggregated values
- **PostgreSQL:** Superior for indexed filtering and complex joins

2. Query Complexity

- **Winner:** MongoDB
- **Reason:** Simpler queries for nested data structures (complexity 1 vs 2-3)

3. Schema Flexibility

- **Winner:** MongoDB
- **Reason:** Dynamic schema, no migrations required

4. Scalability Patterns

- **MongoDB:** Horizontal scaling, document-based sharding

- **PostgreSQL:** Vertical scaling, mature replication

Conclusions and Recommendations

When to Choose MongoDB:

- ✓ **Rapid development** - Flexible schema accelerates prototyping
- ✓ **Document-oriented data** - Products with varying attributes
- ✓ **Read-heavy workloads** - Pre-calculated aggregations
- ✓ **Simple queries** - Embedded data reduces complexity
- ✓ **Horizontal scaling** - Built-in sharding capabilities

When to Choose PostgreSQL:

- ✓ **Complex analytics** - Superior aggregation performance
- ✓ **Strong consistency** - ACID compliance guaranteed
- ✓ **Efficient filtering** - Excellent index performance
- ✓ **Mature ecosystem** - Extensive tooling and expertise

For This E-commerce Use Case:

Recommendation: MongoDB for the following reasons:

1. **Development Speed:** Schema flexibility crucial for evolving product catalogs
2. **Natural Data Model:** Products with embedded reviews match business logic
3. **Read Performance:** Faster for common operations (product browsing, ratings)
4. **Scalability:** Better suited for high-traffic e-commerce platforms

However, consider PostgreSQL if: - Complex reporting and analytics are primary requirements - Team has strong SQL expertise

Personal Reflection

Database Experience

I have Basic knowledge of MongoDB, mainly only from Grundlagen Datenbanken. and good knowledge of SQL from developing web application backends.

Key Learning

The most important insight was seeing how data modeling decisions impact performance at scale. MongoDB's denormalized approach trades storage for query simplicity, while PostgreSQL's normalization optimizes for consistency and complex operations.

Performance vs Complexity Trade-off

MongoDB's simpler queries (complexity 1) can perform better for common operations, but PostgreSQL's sophisticated query optimizer excels at complex analytical workloads.